

What it is, and how you use it.

Type a fashion brand. Get a sourced, one-page opportunity brief — and a real say in it before it counts.

In one breath

You give it a brand name. It **digs the live web** — the brand's site, its product pages, its imagery, its history, its parent group — and writes a one-page brief that answers the seven questions you'd ask before any outreach. The brief ends with a **single wedge** ("here's the gap we'd fill"), a **strategy call** (innovate vs. cut cost), a **ready-to-send opening**, and a **quantified € opportunity** — not just the brand, but its parent group and wholesale partners too.

Two promises hold everywhere: **nothing is fabricated** (every line is sourced, or flagged as an estimate with its assumption), and **the AI never decides alone** — it proposes, you confirm.

HOW IT WORKS

Three layers — and one hard rule

Think of it as a smart researcher who hands you a draft, a calculator the researcher isn't allowed to touch, and you — the operator — with the final word.

1 — AGENTIC CORE • Claude digs live

It runs a tool-use research loop and proposes every field:

web search • catalog/sitemap • PDP audit • product-image
vision • wayback history • people enrichment

Every claim carries a source — or it's flagged as an estimate.

supplies sourced inputs — never the €

2 — DETERMINISTIC SHELL • `economics.py`

€ opportunity — a *band*, not a guess

bottom-up volume × tier cost, cross-checked top-down vs. revenue
+ a 0–5 opportunity score

THE TRUST BOUNDARY — the AI feeds inputs; only this code writes the number.

renders the one-pager → you review

3 — HUMAN LAYER • you

Confirm or sharpen each flagged field, then publish.

The brief shows the AI's draft ↻ and your override ⌘ side by side ⌘.

Decision-makers and contacts always wait for your sign-off.

Your judgment is shown, not assumed — that's where the value is.

The AI proposes • the math is deterministic • you decide

WHY IT'S BUILT THIS WAY

The euro figure and the score are computed by plain, testable Python — the agent can only hand it inputs. So the number can never be "hallucinated," and the same brief is reproducible. It's the one rule the whole system is designed around.

USING IT

Dig → Review → Publish

You work in a small web app. Type a brand, watch it dig, then confirm the judgment. That's the entire loop.



- 1 **Open the app** at `localhost:8501`. In the sidebar, type *any* brand (seeded or not), pick **live**, and hit **Dig** →.
- 2 **It researches the brand** (~1–3 min) and proposes the brief — you see the tool steps as it goes.
- 3 **Review queue** — each flagged field shows the AI's draft and *why*; you confirm it or rewrite it, plus the outreach play (hook / channel / who / opening message).
- 4 **Approve all & publish** — the brief is re-rendered with `✓ confirmed` marks and joins the ranked dashboard.

A live dig costs roughly **\$1** in API credit. With no key, everything still runs in a deterministic offline mode for the three seeded brands.

The seven questions

Every brief answers these — each line sourced, or flagged for you.

Q1

Who is this brand?

Positioning, tier, parent group, markets, channel mix.

Q2

Where is it heading?

Then-vs-now imagery → innovate, or copy what works.

Q3

Momentum?

Growing, pressured — or *split* (brand vs. group diverge).

Q4

How much content do they need?

Images/year across site, marketplace, localized markets.

Q5

What do the product pages look like?

Shots per SKU, on-model vs still-life, consistency — measured.

Q6

Who's the decision-maker?

The role that owns imagery — always flagged for your check.

Q7

How do we get in?

The angle + a sendable opening. Warm path only if evidenced.

+ €


The opportunity

A defensible band for the brand *and* its partners, with the wedge and the strategy call.

WHY YOU CAN TRUST THE NUMBER

How the € is computed

The headline is **Sartiq's** opportunity — not the studio's. We estimate what the brand spends on imagery today (two independent ways, cross-checked), then take the **~20%** Sartiq can realistically capture at its price. Always a range, never a single confident-looking guess — "with this brand, you could do up to this a year."




Two more things make it defensible: it's **tier-aware** (a value retailer isn't priced like a luxury house), and it's **partner-inclusive** — it sizes the whole group and breaks it out *one line per brand* (Diesel · Margiela · Marni ...), because that's where the real opportunity lives.

YOUR PART

Where the human adds value

The machine is good at gathering and proposing. It's deliberately *not* trusted with the calls that need judgment — and those are surfaced to you, not buried:



Example: on a brand whose *group* is shrinking but whose *own* numbers are at a ten-year high, the machine drafts "struggling." You flip it to a **split** read — innovation pitch to the brand, efficiency pitch to the group. The brief shows both, side by side, so the reasoning is visible.

TWO WAYS TO RUN IT

The app, or the terminal

Same engine underneath — pick whichever fits.

The review app · RECOMMENDED

Type a brand, review, publish — no terminal.

```
make setup
source .venv/bin/activate
streamlit run app.py
# → http://localhost:8501
```

The CLI · HEADLESS

Same result, scriptable, good for batches.

```
python3 cli.py "Mango" --live
# opens the brief + dashboard
# --all runs the three seeded brands
# (no key → offline, still works)
```

THE HONEST FINE PRINT

What to keep in mind

The system is built to *tell you* what it doesn't know rather than paper over it:

- **SKU counts and the €** are estimates — shown as bands with stated assumptions, gated by the bottom-up/top-down sanity check.
- **Contacts need validation.** The decision-maker and the way in are always flagged `needs human` — never a warm intro asserted without evidence, never scraped.
- **Live vs. offline is never faked.** A run that falls back is loudly stamped, never dressed up as a live result.